

# Report Monotonicity Extraction

Yichen Han

2025-01-20

```
library(AER)
library(haven)
library(tidyverse)
library(ddml)
library(ranger)
```

## Card (1995)

Adapted analysis plan.

- $N = 3010$  individuals
- $Y$ : continuous, log hourly wage
- $T$ : binary, received education  $> 12$  yrs
- $Z$ : binary, presence of college nearby
- $X$ : 4 binary covariates: work experience  $> 8$  yrs, 3 indicators for race and location of residence.

```
card <- read_dta("https://raw.githubusercontent.com/scunning1975/mixtape/master/card.dta")
suppressMessages(attach(card))
Y <- lwage
D <- as.numeric(educ > 13)
expbin <- as.numeric(exper > 8)
X <- cbind(expbin, black, south, smsa)
Z <- nearc4
idx_complete <- complete.cases(Y, D, X, Z)
Y <- Y[idx_complete]
D <- D[idx_complete]
X <- X[idx_complete, ]
Z <- Z[idx_complete]
df <- data.frame(Y, D, Z, X)
```

If saturated, the data is:

```
X1 <- X[, 1]
X2 <- X[, 2]
X3 <- X[, 3]
X4 <- X[, 4]
X_full <- cbind(X, X1 * X2, X1 * X3, X1 * X4, X2 * X3, X2 * X4, X3 * X4,
               X1 * X2 * X3, X1 * X2 * X4, X1 * X3 * X4, X2 * X3 * X4,
               X1 * X2 * X3 * X4)
Z_int <- cbind(Z, Z * X_full)
```

## DDML

For simplicity, we omit `ivreg`, and only report outputs from DDML.

```
### Estimation DDML
# ----- estimating beta_rich using DDML
set.seed(123)
# Estimate the local average treatment effect using short-stacking with base
#   learners ols, rlasso, and xgboost.
learners_multiple <- list(list(fun = ols),
                           list(fun = mdl_glmnet),
                           list(fun = mdl_ranger),
                           list(fun = mdl_xgboost))
pliv_rich <- ddml_pliv(Y, D, Z, X,
                      learners = learners_multiple,
                      ensemble_type = c('nnls', 'nnls1', 'ols', 'average'),
                      shortstack = TRUE,
                      sample_folds = 10,
                      silent = TRUE)

summary(pliv_rich)
```

```
## PLIV estimation results:
##
## , , nnls
##
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.00112    0.00928   0.121  0.9040
## D_r          0.91380    0.43327   2.109  0.0349
##
## , , nnls1
##
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  5.12e-05    0.00926  0.00553  0.996
## D_r          9.14e-01    0.43338  2.10889  0.035
##
## , , ols
##
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.000977    0.00926   0.105  0.9160
## D_r          0.912894    0.43388   2.104  0.0354
##
## , , average
##
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  7.99e-05    0.00924  0.00865  0.9931
## D_r          9.08e-01    0.43141  2.10440  0.0353
```

```
# ----- estimating LATE using DDML
late_fit_short <- ddml_late(Y, D, Z, X,
                           learners = learners_multiple,
                           ensemble_type = c('nnls', 'nnls1', 'ols', 'average'),
                           shortstack = TRUE,
                           sample_folds = 10,
```

```

                                silent = TRUE)
summary(late_fit_short)

```

```

## LATE estimation results:
##
##      Estimate Std. Error t value Pr(>|t|)
## nnls      0.670     0.323   2.08  0.0378
## nnls1     0.671     0.322   2.08  0.0374
## ols       0.696     0.330   2.11  0.0349
## average   0.668     0.320   2.09  0.0367

```

## Weight Estimation

Note that DDML-PLIV outputs  $\beta_{\text{rich}}$ , and for the binary case

$$\beta_{\text{rich}} = \mathbb{E} \left[ \beta_{\text{acr}}(X) \frac{P[T(1) > T(0)|X] \mathbb{V}[Z|X]}{\mathbb{E}[P[T(1) > T(0)|X] \mathbb{V}[Z|X]]} \right]$$

where  $\beta_{\text{acr}}(X) := \mathbb{E}[Y(T(1)) - Y(T(0)) | T(1) > T(0)]$ .

We regard  $\frac{P[T(1) > T(0)|X] \mathbb{V}[Z|X]}{\mathbb{E}[P[T(1) > T(0)|X] \mathbb{V}[Z|X]]}$  as the weights.

We estimate  $P[T(1) > T(0)|X]$  and  $\mathbb{V}[Z|X]$  using two random forests.

First, define the function to estimate weights.

```

estimate_rich_weights <- function(df, z_vars = "Z", plot=TRUE) {
  set.seed(123)
  # 1) Identify covariates: all columns except Y, D, Z
  covariates <- setdiff(names(df), c("Y", "D", z_vars))

  # 2) Ensure D, Z are factors (binary for this example)
  df$D <- as.factor(df$D)
  df$Z <- as.factor(df$Z)

  # 3) Fit model for D ~ Z + X
  formula_D <- as.formula(
    paste("D ~ ", paste(c(z_vars, covariates), collapse = " + "))
  )
  model_D <- ranger(formula_D, data = df, probability = TRUE)

  # 4) Predict P(D=1) for Z=1 vs Z=0 to get p_diff = P[T(1)>T(0)|X]
  p1_hat <- predict(model_D, data = transform(df, Z=1), type = "response")$predictions[,2]
  p0_hat <- predict(model_D, data = transform(df, Z=0), type = "response")$predictions[,2]
  p_diff <- p1_hat - p0_hat # estimated compliance propensity

  # 5) Fit model for Z ~ X to get pZ_hat, then var(Z|X) = pZ_hat*(1 - pZ_hat)
  formula_Z <- as.formula(
    paste("Z ~", paste(covariates, collapse = " + "))
  )
  model_Z <- ranger(formula_Z, data = df, probability = TRUE)

  pZ_hat <- predict(model_Z, data = df, type = "response")$predictions[,2]
}

```

```

vZ_hat <- pZ_hat * (1 - pZ_hat)

# 6) Final weights = p_diff * var(Z|X), normalized to sum to 1
w_raw <- p_diff * vZ_hat
w_final <- w_raw / sum(w_raw)
df$w_final <- w_final

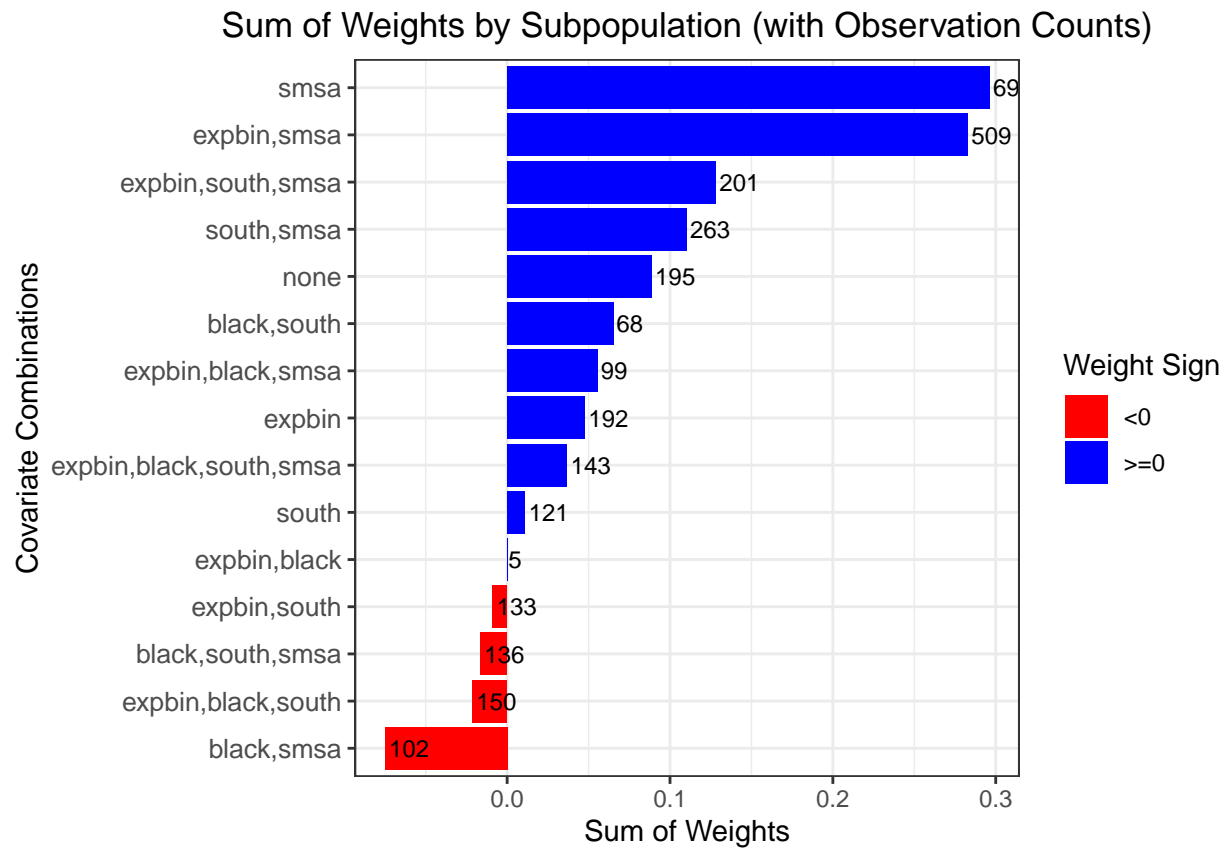
# 7) Create a 'group' label for every unique combination of X
df$group <- apply(df[, covariates, drop = FALSE], 1, function(xx) {
  # Exclude variables matching the pattern "V<number>"
  vars_to_include <- names(xx)[xx == 1 & !grepl("^V[0-9]+.*$", names(xx))]
  paste(vars_to_include, collapse = ",")
})
df$group[df$group == ""] <- "none" # Label groups with no variables set to 1

# 8) Sum weights by group
group_summary <- aggregate(w_final ~ group, data = df, sum)
group_summary$obs_count <- aggregate(w_final ~ group, data = df, length)$w_final # Count observation
group_summary$sign <- ifelse(group_summary$w_final >= 0, ">=0", "<0") # Positive/negative sign


if (plot) {
  pt <- ggplot(group_summary, aes(x = reorder(group, w_final), y = w_final, fill = sign)) +
    geom_bar(stat = "identity") +
    geom_text(
      aes(label = obs_count),
      hjust = -0.1, # Position labels slightly to the right of the bars
      size = 3
    ) +
    coord_flip() + # Flip axes for better readability
    scale_fill_manual(values = c(">=0" = "blue", "<0" = "red")) +
    labs(
      title = "Sum of Weights by Subpopulation (with Observation Counts)",
      x = "Covariate Combinations",
      y = "Sum of Weights",
      fill = "Weight Sign"
    ) +
    theme_bw() +
    theme(
      axis.text.y = element_text(size = 10),
      plot.title = element_text(hjust = 0.5)
    )
}
print(pt)
# Return a list with the augmented df and group-weight table
return(list(
  df_with_weights = df,
  group_summary = group_summary
))
}

```

Estimate weights under saturated  $X$  and simple  $Z$ .



Estimate weights under saturated  $X$  and saturated  $Z$  (Saturate and Weight).



## Monotonicity Extraction

Looking at the graphs, we suspect, under rich covariates,

- black, smsa, not south, not expbin
- expbin, black, south, not smsa
- expbin, south, not smsa, not black
- black, south, smsa, not expbin

are not correct in monotonicity.

Further, under Saturation and Weight,

- black, south, smsa, not expbin
- black, smsa, not south, not expbin

are with some confidence, defiers.

Hypothesis: removing them from the data and doing PLIV again will yield LATE.

First remove all four groups.

```
df_comp <- df %>%
  filter(!(black == 1 & smsa == 1 & expbin == 0 & south == 0),
         !(expbin == 1 & black == 1 & south == 1 & smsa == 0),
         !(expbin == 1 & south == 1 & black == 0 & smsa == 0),
```

```

      !(black == 1 & south == 1 & smsa == 1 & expbin == 0))
Y <- df_comp$Y
D <- df_comp$D
Z <- df_comp$Z
X_comp <- df_comp %>% select(-Y, -D, -Z) %>% as.matrix()
# redo PLIV
set.seed(123)
pliv_comp <- ddml_pliv(Y, D, Z, X_comp,
                      learners = learners_multiple,
                      ensemble_type = c('nnls', 'nnls1', 'ols', 'average'),
                      shortstack = TRUE,
                      sample_folds = 10,
                      silent = TRUE)
summary(pliv_comp)

```

```

## PLIV estimation results:
##
## , , nnls
##
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.000626   0.00903  -0.0693   0.9448
## D_r          0.651820   0.28291   2.3040   0.0212
##
## , , nnls1
##
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.000218   0.00903  -0.0241   0.9808
## D_r          0.651804   0.28291   2.3039   0.0212
##
## , , ols
##
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.000752   0.00902   0.0834   0.9335
## D_r          0.646825   0.27962   2.3133   0.0207
##
## , , average
##
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.000223   0.00899  -0.0249   0.9802
## D_r          0.639203   0.28068   2.2773   0.0228

```

Clearly, the hypothesis is supported by the data.

Then remove the two groups with some confidence.

```

df_comp <- df %>%
  filter(!(black == 1 & smsa == 1 & expbin == 0 & south == 0),
         !(black == 1 & south == 1 & smsa == 1 & expbin == 0))
Y <- df_comp$Y
D <- df_comp$D
Z <- df_comp$Z
X_comp <- df_comp %>% select(-Y, -D, -Z) %>% as.matrix()
# redo PLIV

```

```

set.seed(123)
pliv_comp <- ddml_pliv(Y, D, Z, X_comp,
                      learners = learners_multiple,
                      ensemble_type = c('nnls', 'nnls1', 'ols', 'average'),
                      shortstack = TRUE,
                      sample_folds = 10,
                      silent = TRUE)
summary(pliv_comp)

```

```

## PLIV estimation results:
##
## , , nnls
##
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.000466   0.00908  -0.0513   0.9591
## D_r          0.800774   0.33954   2.3584   0.0184
##
## , , nnls1
##
##           Estimate Std. Error  t value Pr(>|t|)
## (Intercept) -1.63e-05   0.00908  -0.00179   0.9986
## D_r          8.01e-01   0.33956   2.35838   0.0184
##
## , , ols
##
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) 0.000279   0.00906   0.0308   0.975
## D_r          0.795931   0.33632   2.3666   0.018
##
## , , average
##
##           Estimate Std. Error  t value Pr(>|t|)
## (Intercept) -4.57e-05   0.00904  -0.00506   0.9960
## D_r          7.89e-01   0.33797   2.33546   0.0195

```

The result is not entirely the same, so removing only the best-suspected defiers is not enough.